

Trishal Muthan

Period 2

Outstanding Work: Book Report

You Look Like a Thing and I Love You: How Artificial Intelligence Works and Why It's Making the World a Weirder Place: by Janelle Shane

You Look Like a Thing and I Love You: How Artificial Intelligence Works and Why It's Making the World a Weirder Place is a particularly interesting book. It's actually the first nonfiction book I have ever read about artificial intelligence and I enjoyed it thoroughly. It was a very new and unusual experience to be reading about something that we are working on every week and something that is evolving at such a rapid pace. Topics which the book covered, ranging from things like neural networks and random forests (common AI techniques) to bias and shortcuts (common errors that AI can encounter), are things I will likely work with during my time at TJ and in the future. As a result, books like this not only allow me to get away from the computer screen but also allow me to learn more about subjects that will be especially useful down the road.

I think a common misconception that most of society has nowadays is that AI is something to be wary or nervous of. And until quite recently, it's a misconception that I've held myself. In addition to the good chunk of people who perceive AI as something that might wage war on humans (like something out of a science-fiction novel or movie), there are also those who think that AI will soon be taking over jobs and society as a whole. The book places a huge emphasis on making the true expectations we should have for AI in the near future especially clear to the reader. The reality of the matter is that AI is not even close to what we see in fiction or what many think of when they think of AI. Not in the slightest. This is extremely relevant nowadays as people begin relying more on automation and AI to do things that are too tedious or unnecessary for humans. Shane describes many of the different limitations to AI that we currently face and will have to overcome if we want to get even remotely close to any current expectations we have for AI. I found learning about these limitations very surprising because I never realized how much is left for us to fix if we want to really see drastic improvements on AI technology. Not only are there many limitations, but these limitations are in a variety of different areas. For example, one limitation of AI is the development of bias while training. Shane gives an instance of this in an AI powered hiring process: "If the algorithm is trained in the way that human hiring managers have selected or ranked resumes in the past, it's very likely to pick up bias." Putting an AI that has the ability to develop a bias into use is very easy and so a lot of effort must be put into ensuring that the AI is fair. Other limitations of AI include insufficient or errors in data, incorrect training, or a flawed application of AI. All of these different limitations that Shane described in the book really showed me that there is significantly more work that needs to be done in the field of AI if we are going to get anywhere in pushing the extent of capability that AI can potentially have.

This book has honestly made my perspective on the field of computer science and AI more realistic as a whole. Shane did a great job in conveying the true nature of the field and where it currently stands to the reader and I think that my viewpoint on AI has changed considerably. I think that I was expecting a bit too much out of the field at this point at this point and time. I thought that the field of AI was more advanced and closer to being some of the things that many people are wary of than it really is. And while the field is still extremely extensive and is growing at a faster and faster pace, it is not where I thought it had been. But all that this means is that there are clear problems that researchers, engineers, and maybe even myself are able to focus on.